

Support Vector Machines

SYS 6018 | Spring 2023

svm.pdf

Contents

1	Support Vector Machines (SVM) Introduction	2
1.1	Required R Packages	2
1.2	Example	2
1.3	SVM as Loss + Penalty	3
1.4	SVM vs. Logistic Regression	4
1.5	Example: Compare Linear Classifiers	5
1.6	Three Representations of the Optimization Problem	6
1.7	Linear SVM	6
2	Kernels and Non-linear SVM	7
2.1	Basis Expansion	7
2.2	Alternative (Dual) Formulation	8
2.3	SVM in R	8
2.4	Kernels	9
2.5	Unbalanced Data	12
3	Appendix: R Code	13

1 Support Vector Machines (SVM) Introduction


1.1 Required R Packages

We will be using the R packages of:

- `tidyverse` for data manipulation and visualization
- `e1071` for the `svm()` functions

1.2 Example

Goal: find *best* line(s)/curve(s) to separate the two classes.



1.3 SVM as Loss + Penalty

Training Data: $\{(\tilde{y}_i, \mathbf{x}_i)\}_{i=1}^n$

- $\tilde{y}_i \in \{-1, +1\}$
- $\mathbf{x}_i^\top = [x_{i1}, x_{i2}, \dots, x_{in}] \in \mathbb{R}^p$

Predictor Function: $f(x)$


- Linear: $f(\mathbf{x}; \beta) = \beta_0 + \sum_{j=1}^p x_j \beta_j$
- Basis Expansion: $f(\mathbf{x}; \beta) = \beta_0 + \sum_{j=1}^d h_j(\mathbf{x}) \beta_j$
 - $h_j(\mathbf{x})$ transforms the raw x vector
 - Polynomial example: $h_1(\mathbf{x}) = x_1$, $h_2(\mathbf{x}) = x_1^2$, $h_3(\mathbf{x}) = x_2$, $h_4(\mathbf{x}) = x_2^2$, $h_5(\mathbf{x}) = x_1 x_2$
 - This is the connection to *kernels* that we will discuss later
- Let $f_i = f(\mathbf{x}_i)$

Classification:

- If $\hat{f}_i \geq 0$, then label as class +1
- If $\hat{f}_i < 0$, then label as class -1

Loss Function: Hinge Loss

$$L(\tilde{y}_i, f_i) = \max\{0, 1 - \tilde{y}_i f_i\}$$



Penalty Function: Ridge Penalty

$$\begin{aligned} P_\lambda(\beta) &= \frac{\lambda}{2} \sum_{j=1}^d \beta_j^2 \\ &= \lambda \|\beta\|^2 / 2 \end{aligned}$$

- This should have you thinking that the \mathbf{x} 's should be scaled!

Summary of SVM

1. Estimate *model parameters* β

$$\begin{aligned}\hat{\beta}_\lambda &= \arg \min_{\beta} \left\{ \sum_{i=1}^n \max\{0, 1 - \tilde{y}_i f_i(\beta)\} + \lambda \|\beta\|^2 / 2 \right\} \\ &= \arg \min_{\beta} \{ \text{Hinge Loss}(\beta) + \lambda \text{Penalty}(\beta) \}\end{aligned}$$

2. Label Class +1 if $\hat{f}_i \geq 0$, else label Class -1
 - This implicitly assumes a 0-1 loss (or equal cost FP and FN)
 - More generally, use threshold t that considers the costs of FP and FN
3. **Tuning Parameters:** besides the λ , SVM will also have tuning parameters related to the kernels (more to come on this).

1.4 SVM vs. Logistic Regression

The *linear* SVM is actually very similar to Logistic Regression (with Ridge Penalty).

Let $f_i(\beta) = \beta_0 + \sum_{j=1}^p x_j \beta_j$ be the log-odds.

- We also referred to this as $\gamma(\mathbf{x}_i)$ in previous notes.


1. Estimate *model parameters* β
 - This assumes standardized \mathbf{x} 's

$$\begin{aligned}\hat{\beta}_\lambda &= \arg \min_{\beta} \left\{ \sum_{i=1}^n \log(1 + \exp(-\tilde{y}_i f_i(\beta))) + \lambda \|\beta\|^2 / 2 \right\} \\ &= \arg \min_{\beta} \{ \text{Log Loss}(\beta) + \lambda \text{Penalty}(\beta) \}\end{aligned}$$

2. Label Class +1 if $\hat{f}_i \geq t$, else label Class -1
 - For some threshold t that considers the costs of FP and FN
 - If $\text{Cost}(\text{FP}) = \text{Cost}(\text{FN})$, the set $t = 0$, which is equivalent to $p(x) = 1/2$


Details of the Logistic Regression Loss Function

Note



The Log-Loss has been scaled so it equals the Hinge Loss at $f=0$.

1.5 Example: Compare Linear Classifiers



1.6 Three Representations of the Optimization Problem

1. Penalized optimization

$$\begin{aligned}\hat{\beta} &= \arg \min_{\beta} \{ \text{Loss}(\beta) + \lambda \text{Penalty}(\beta) \} \\ &= \arg \min_{\beta} \{ C \text{Loss}(\beta) + \text{Penalty}(\beta) \}\end{aligned}$$

- where $C = \frac{1}{\lambda} > 0$ is an alternative strength of penalty
- In SVM, C is referred to as the *cost*

2. Constraint on Penalty

$$\begin{aligned}\hat{\beta} &= \arg \min_{\beta} \text{Loss}(\beta) \quad \text{subject to } \text{Penalty}(\beta) \leq t \\ &= \arg \min_{\beta: \text{Penalty}(\beta) \leq t} \text{Loss}(\beta)\end{aligned}$$


3. Constraint on Loss

$$\begin{aligned}\hat{\beta} &= \arg \min_{\beta} \text{Penalty}(\beta) \quad \text{subject to } \text{Loss}(\beta) \leq M \\ &= \arg \min_{\beta: \text{Loss}(\beta) \leq M} \text{Penalty}(\beta)\end{aligned}$$

Note

There is a one-to-one relationship between all tuning parameters: λ , C , t , M .

1.7 Linear SVM






2 Kernels and Non-linear SVM

2.1 Basis Expansion

- Linear: $f(\mathbf{x}; \beta) = \beta_0 + \sum_{j=1}^p x_j \beta_j$
- Basis Expansion: $f(\mathbf{x}; \beta) = \beta_0 + \sum_{j=1}^d h_j(\mathbf{x}) \beta_j$
 - $h_j(\mathbf{x})$ transforms the raw \mathbf{x} vector

2.1.1 Polynomial Expansion: Quadratic Terms

- Polynomial example: $h_1(\mathbf{x}) = x_1$, $h_2(\mathbf{x}) = x_1^2$, $h_3(\mathbf{x}) = x_2$, $h_4(\mathbf{x}) = x_2^2$, $h_5(\mathbf{x}) = x_1 x_2$
- $f(\mathbf{x}; \beta) = \beta_0 + \sum_{j=1}^5 h_j(\mathbf{x}) \beta_j$



2.2 Alternative (Dual) Formulation

It turns out that the SVM solution for the model coefficients β can be written as

$$\hat{\beta}_j = \sum_{i=1}^n \hat{\alpha}_i \tilde{y}_i h_j(\mathbf{x}_i)$$

- See ESL Eq. 12.17 if you are interested in the details.
- In this reformulation, the $\{\hat{\alpha}_i\}_{i=1}^n$ become the model parameters.
 - There is one model parameter for each observation!
 - $0 \leq \alpha_i \leq C$ (or $0 \leq \alpha_i \leq \frac{1}{\lambda}$)
 - But SVM will force most values to 0.
 - The observations with $\alpha_i > 0$ are called the *support vectors*
 - * So the entire SVM model is a function of the support vectors only!
 - * The support vectors are the observations on the wrong side of the margin.
- The decision function $\hat{f}(\mathbf{x})$ can be re-written

$$\begin{aligned} \hat{f}(\mathbf{x}) &= f(\mathbf{x}; \hat{\beta}) \\ &= \hat{\beta}_0 + \sum_{j=1}^d h_j(\mathbf{x}) \hat{\beta}_j \\ &= \hat{\beta}_0 + \sum_{j=1}^d h_j(\mathbf{x}) \left[\sum_{i=1}^n \hat{\alpha}_i \tilde{y}_i h_j(\mathbf{x}_i) \right] \\ &= \hat{\beta}_0 + \sum_{i=1}^n \hat{\alpha}_i \tilde{y}_i \left[\sum_{j=1}^d h_j(\mathbf{x}) h_j(\mathbf{x}_i) \right] \\ &= \hat{\beta}_0 + \sum_{i=1}^n \hat{\alpha}_i \tilde{y}_i K(\mathbf{x}, \mathbf{x}_i) \end{aligned}$$

where

$$K(\mathbf{x}, \mathbf{x}_i) = \sum_{j=1}^d h_j(\mathbf{x}) h_j(\mathbf{x}_i) = \langle h(\mathbf{x}), h(\mathbf{x}_i) \rangle$$


is called a **kernel** and measures the inner product, or *similarity* between x and x_i (observation i).

2.3 SVM in R

- The `svm()` function from the `e1071` package can implement SVM.
 - There is also a helpful `tune.svm()` to help you select the tuning parameters.
- The ISLR Lab in Section 9.6 has example R code.

2.4 Kernels

To help illustrate the difference between different kernels, let's look at slightly different data that won't be easy to classify using linear classifiers.




Three popular kernels (but there are many more) are the *linear*, *polynomial*, and *radial basis*

2.4.1 Linear Kernel

The *linear kernel* is

$$K(\mathbf{x}, \mathbf{u}) = \sum_{j=1}^p \mathbf{x}_j \mathbf{u}_j = \langle \mathbf{x}, \mathbf{u} \rangle$$




$$\begin{aligned} \hat{f}_{\text{linear}}(\mathbf{u}) &= \hat{\beta}_0 + \sum_{i=1}^n \hat{\alpha}_i \tilde{y}_i K(\mathbf{x}_i, \mathbf{u}) \\ &= \hat{\beta}_0 + \sum_{i=1}^n \hat{\alpha}_i \tilde{y}_i \left(\sum_{j=1}^p \mathbf{x}_{ij} \mathbf{u}_j \right) \end{aligned}$$

2.4.2 Polynomial Kernel

The *polynomial kernel* of degree deg is

$$K(\mathbf{x}, \mathbf{u}) = \left(1 + \sum_{j=1}^p \mathbf{x}_j \mathbf{u}_j \right)^{\text{deg}}$$

Note: in R, the `'svm()'` function from `'e1071'` package includes two other tuning parameters (γ and `coef0`)



$$\begin{aligned} \hat{f}_{\text{poly}}(\mathbf{u}) &= \hat{\beta}_0 + \sum_{i=1}^n \hat{\alpha}_i y_i K(\mathbf{x}_i, \mathbf{u}) \\ &= \hat{\beta}_0 + \sum_{i=1}^n \hat{\alpha}_i y_i \left(1 + \sum_{j=1}^p \mathbf{x}_j \mathbf{u}_j \right)^{\text{deg}} \end{aligned}$$


2.4.3 Radial Basis Kernel




The *Radial Basis Kernel* with parameter γ (γ is gamma) is

$$\begin{aligned} K(\mathbf{x}, \mathbf{u}) &= \exp \left(-\gamma \sum_{j=1}^p (\mathbf{x}_j - \mathbf{u}_j)^2 \right) \\ &= \exp \left(-\gamma \text{dist}^2(\mathbf{x}, \mathbf{u}) \right) \end{aligned}$$

where $\text{dist}^2(\mathbf{x}, \mathbf{u})$ is the *squared Euclidean distance* between \mathbf{x} and \mathbf{u} .

- The Radial Basis kernel is large for test observations close to training observations.
- Notice that both `cost` and `gamma` are influential tuning parameters.





$$\begin{aligned}\hat{f}_{\text{radial}}(\mathbf{u}) &= \hat{\beta}_0 + \sum_{i=1}^n \hat{\alpha}_i \tilde{y}_i K(\mathbf{x}_i, \mathbf{u}) \\ &= \hat{\beta}_0 + \sum_{i=1}^n \hat{\alpha}_i \tilde{y}_i \exp\left(-\gamma \text{dist}^2(\mathbf{x}_i, \mathbf{u})\right)\end{aligned}$$

2.5 Unbalanced Data

Three primary approaches:

1. Use weights corresponding to prior class probabilities
 - See the `class.weights` argument in `?svm`
 - Ensure the resulting score is what you want for unadjusted prediction.

$$\hat{\beta}_{\lambda} = \arg \min_{\beta} \left\{ \sum_{i=1}^n w_i \max\{0, 1 - \tilde{y}_i f_i(\beta)\} + \lambda \|\beta\|^2 / 2 \right\}$$

2. Use a threshold on the score
 - E.g., choose label = +1 if: $\hat{f}(\mathbf{x}) > t$
 - The \hat{f} are the `decision.values` output of `predict.svm`
 - Its an attribute of the returned object
3. Sampling
 - under or over sample training data from one class to balance the label distribution.
 - The SVM (hinge) objective function is optimized at $\text{sign}[\Pr(Y = +1|X = x) - 1/2]$.
 - Ensure the resulting score is what you want for unadjusted prediction.

3 Appendix: R Code


```
#-----#
# Create Data
#-----#
library(mvtnorm)

prior = c(.60, .40)
mu1 = c(2, 2)
mu2 = c(0, 0)
mu2.B = c(4, 4)
sigma1 = .5*matrix(c(2,-1, -1, 2), nrow=2)
sigma2 = .5*matrix(c(1,0,0,1), nrow=2)

set.seed(2020)
n = 200
n1 = rbinom(1, size=n, prob=prior[1])
n2 = n-n1
n2 = round(c(n2/2, n2/2))
X1 = mvtnorm::rmvnorm(n1, mean=mu1, sigma = sigma1)
X2 = mvtnorm::rmvnorm(n2[1], mean=mu2, sigma = sigma2)
X2 = rbind(X2, mvtnorm::rmvnorm(n2[2], mean=mu2.B, sigma = sigma2))
labels = c("+1", "-1")

data_mix = bind_rows(
  !!labels[1] := as_tibble(X1, .name_repair = ~str_c("X", seq_along(.))),
  !!labels[2] := as_tibble(X2, .name_repair = ~str_c("X", seq_along(.))),
  .id = "class"
)

data_mix %>%
  ggplot(aes(X1, X2, color = class)) +
  geom_point()
```



```
#-----#
# Create Cross-Validation Folds
#-----#
set.seed(2022)
n.folds = 10
fold = sample(rep(1:n.folds, length=nrow(data_mix)))

# using tidymodels
library(rsample)
set.seed(2022)
cv_folds = rsample::vfold_cv(data_mix, v = 10)

#-----#
# Polynomial SVM Example
#-----#
library(e1071)
fit = svm(factor(class) ~ X1 + X2,
  data = data_mix[fold != 1, ], # hold out fold 1
  #: tuning parameters
  kernel = "poly",
  degree = 2,
  cost = 1
)

# Notes:
# - svm() requires that the outcome variable be a "factor" for classification
#   problems.
# - the polynomial kernel only has cost and degree parameters

#: basic model summary
summary(fit)
#>
#> Call:
#> svm(formula = factor(class) ~ X1 + X2, data = data_mix[fold != 1,
#>   ], kernel = "poly", degree = 2, cost = 1)
#>
#>
```

```

#> Parameters:
#>   SVM-Type:  C-classification
#>   SVM-Kernel: polynomial
#>     cost:  1
#>   degree:  2
#>   coef.0:  0
#>
#> Number of Support Vectors:  32
#>
#> ( 17 15 )
#>
#>
#> Number of Classes:  2
#>
#> Levels:
#>  -1 +1

# : predictions
pred = predict(fit, data_mix[fold == 1,], decision.values = TRUE)

eval_data = tibble(
  outcome = data_mix$class[fold == 1],
  pred_hard = c(pred),
  pred_soft = as.numeric(attr(pred, "decision.values"))
)

# Notes:
# - using the `decision.values = TRUE` argument to get the scores. But set
#   to default of FALSE to just get the hard classifications
# - the scores can be used to create rank or threshold-based metrics like ROC
#   curves
# - There is also a way to get probability estimates by setting
#   svm(..., probability = TRUE) and predict(..., probability = TRUE).
#   This attempts to convert scores to probabilities

# : evaluation
library(yardstick)
levs = c("+1", "-1") # set outcome of interest at first level

eval_data %>%
  yardstick::roc_curve(truth = factor(outcome, levels=levs), estimate = pred_soft)
#> # A tibble: 22 x 3
#>   .threshold specificity sensitivity
#>   <dbl>         <dbl>         <dbl>
#> 1    -Inf             0             1
#> 2   -11.9             0             1
#> 3    -7.08            0.1             1
#> 4    -4.82            0.2             1
#> 5    -4.13            0.3             1
#> 6    -3.40            0.4             1
#> # ... with 16 more rows

eval_data %>%
  yardstick::accuracy(truth = factor(outcome, levels=levs),
    estimate = factor(pred_hard, levels=levs))
#> # A tibble: 1 x 3
#>   .metric .estimator .estimate
#>   <chr>   <chr>         <dbl>

```

```

#> 1 accuracy binary          0.95

# Notes:
# - the yardstick packages wants outcome variables and hard predictions be
#   "factors" with the outcome of interest as the first level

#-----#
# Radial basis SVM tuning parameter selection
#-----#
#:: create function to fit, predict, and evaluate
eval_svm <- function(data_train, data_test, cost, gamma){

  #:: fit
  fit = svm(factor(class) ~ X1 + X2,
            data = data_train,
            #:: tuning parameters
            kernel = "radial",
            gamma = gamma,
            cost = cost
            )

  #:: predict
  pred = predict(fit, data_test, decision.values = TRUE)

  #:: evaluate
  eval_data = tibble(
    outcome = data_test$class,
    pred_hard = c(pred),
    pred_soft = as.numeric(attr(pred, "decision.values"))
  )

  levs = c("+1", "-1") # set outcome of interest at first level

  auroc = eval_data %>%
    yardstick::roc_auc(truth = factor(outcome, levels=levs), estimate = pred_soft)

  accuracy = eval_data %>%
    yardstick::accuracy(truth = factor(outcome, levels=levs),
                       estimate = factor(pred_hard, levels=levs))

  bind_rows(auroc, accuracy) %>%
    mutate(cost, gamma) # add tuning parameters
}

#:: test it out
eval_svm(
  data_train = data_mix[fold != 1,],
  data_test = data_mix[fold == 1,],
  cost = 1,
  gamma = .1
)
#> # A tibble: 2 x 5
#>   .metric .estimator .estimate cost gamma
#>   <chr>    <chr>         <dbl> <dbl> <dbl>
#> 1 roc_auc binary          0.99     1    0.1
#> 2 accuracy binary          0.95     1    0.1

```



```

# using tidymodels
library(rsample)
eval_svm(
  data_train = training(cv_folds$splits[[1]]),
  data_test = testing(cv_folds$splits[[1]]),
  cost = 1,
  gamma = .1
)
#> # A tibble: 2 x 5
#>   .metric .estimator .estimate cost gamma
#>   <chr>    <chr>      <dbl> <dbl> <dbl>
#> 1 roc_auc binary      0.99     1  0.1
#> 2 accuracy binary      0.95     1  0.1

# create grid of tuning parameter values
tune_grid = expand_grid(
  cost = 10^(-2:2),
  gamma = c(.1, .5, 1)
)

# loop over folds and grids
out = tibble()

for(k in unique(fold)){
  for(i in 1:nrow(tune_grid)){
    metrics = eval_svm(
      data_train = data_mix[fold != k,],
      data_test = data_mix[fold == k,],
      cost = tune_grid$cost[i],
      gamma = tune_grid$gamma[i]
    )

    out = bind_rows(out, metrics %>% mutate(fold = k) )
  }
}

# average over folds
out %>%
  # calculate average performance over folds
  group_by(.metric, cost, gamma) %>%
  summarize(mu = mean(.estimate), .groups = "drop") %>%
  # spread data wider to see both metrics
  pivot_wider(names_from = .metric, values_from = mu) %>%
  # arrange by auc and then accuracy (if ties)
  arrange(desc(roc_auc), desc(accuracy))
#> # A tibble: 15 x 4
#>   cost gamma accuracy roc_auc
#>   <dbl> <dbl>    <dbl>    <dbl>
#> 1 100     0.1    0.925    0.984
#> 2  10     0.5    0.935    0.982
#> 3  0.1    0.5    0.935    0.982
#> 4  1      0.1    0.945    0.981
#> 5  10     1      0.93     0.981
#> 6  10     0.1    0.93     0.980

```

```
#> # ... with 9 more rows

# Notes:
# - can go back and modify grid to focus on best tuning region

# using tidymodels and purrr

# function to search over grid
eval_svm_grid <- function(data, grid){
  pmap_df(grid, function(cost, gamma) eval_svm(training(data), testing(data), cost, gamma))
}

# try it out
data = cv_folds$splits[[1]]
eval_svm_grid(data, head(tune_grid))
#> # A tibble: 12 x 5
#>   .metric .estimator .estimate cost gamma
#>   <chr>    <chr>         <dbl> <dbl> <dbl>
#> 1 roc_auc binary         0.97  0.01  0.1
#> 2 accuracy binary         0.5   0.01  0.1
#> 3 roc_auc binary         0.98  0.01  0.5
#> 4 accuracy binary         0.5   0.01  0.5
#> 5 roc_auc binary         0.97  0.01  1
#> 6 accuracy binary         0.5   0.01  1
#> # ... with 6 more rows

# loop over folds
map_df(cv_folds$splits, ~eval_svm_grid(data=.x, grid=tune_grid), .id="fold")
#> # A tibble: 300 x 6
#>   fold .metric .estimator .estimate cost gamma
#>   <chr> <chr>    <chr>         <dbl> <dbl> <dbl>
#> 1 1      roc_auc binary         0.97  0.01  0.1
#> 2 1      accuracy binary         0.5   0.01  0.1
#> 3 1      roc_auc binary         0.98  0.01  0.5
#> 4 1      accuracy binary         0.5   0.01  0.5
#> 5 1      roc_auc binary         0.97  0.01  1
#> 6 1      accuracy binary         0.5   0.01  1
#> # ... with 294 more rows
```